

# Assignment 3

Jacqueline Wu, Sam Zhang

March 2024

## 1 Value-based method with linear function approximation

We implement Q-Learning and Expected SARSA and experiment them with two environments from the Gymnasium, an open-source fork of the OpenAI Gym. [1] During the experiment, the environment's states are discretized through tile-coding into a binary vector.

### 1.1 Results

We experiment with 3  $\epsilon$  per algorithm-environment pair, and use a fixed set of 3 different learning rates  $\alpha \in \{\frac{1}{4}, \frac{1}{8}, \frac{1}{16}\}$ . The results are shown in Figure 1.

### 1.2 Discussion

## 2 Policy Gradient Theorem

## 3 Policy-based methods with linear function approximation

### 3.1 Results

### 3.2 Discussion

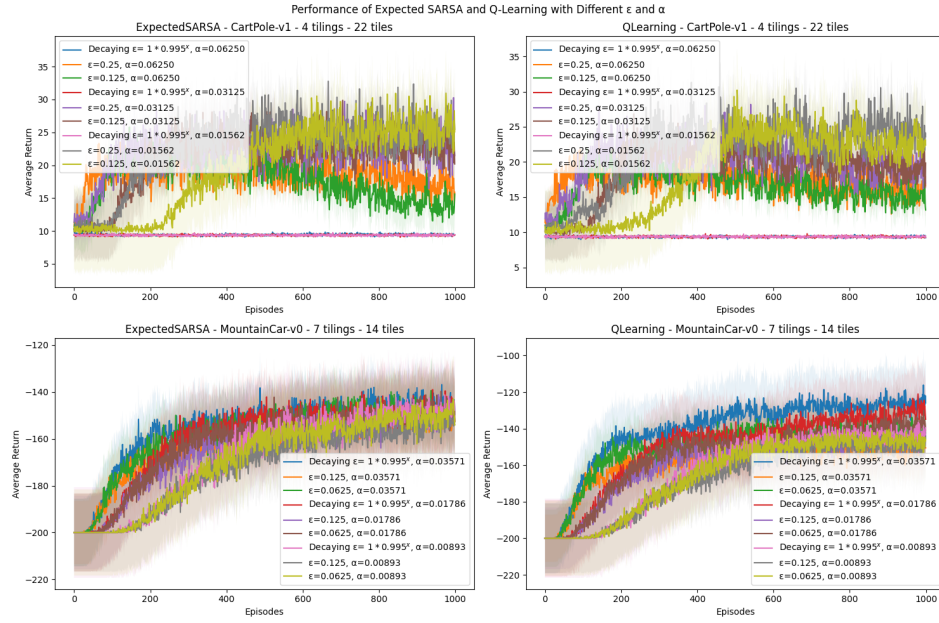


Figure 1: ExpectedSARSA and Q-Learning's Performances on CartPole-v1 and MountainCar-v0

## References

- [1] Mark Towers, Jordan K. Terry, Ariel Kwiatkowski, John U. Balis, Gianluca de Cola, Tristan Deleu, Manuel Goulão, Andreas Kallinteris, Arjun KG, Markus Krimmel, Rodrigo Perez-Vicente, Andrea Pierré, Sander Schulhoff, Jun Jet Tai, Andrew Tan Jin Shen, and Omar G. Younis. Gymnasium, March 2023.